

```
#hide
! [ -e /content ] && pip install -Uqq fastbook
import fastbook
fastbook.setup_book()

_____ 719.8/719.8 KB 10.3 MB/s eta 0:00:00
_____ 6.7/6.7 MB 46.3 MB/s eta 0:00:00
_____ 1.3/1.3 MB 16.9 MB/s eta 0:00:00
_____ 469.0/469.0 KB 7.6 MB/s eta 0:00:00
_____ 132.9/132.9 KB 8.7 MB/s eta 0:00:00
_____ 212.2/212.2 KB 22.0 MB/s eta 0:00:00
_____ 1.0/1.0 MB 28.4 MB/s eta 0:00:00
_____ 199.2/199.2 KB 11.9 MB/s eta 0:00:00
_____ 110.5/110.5 KB 6.8 MB/s eta 0:00:00
_____ 7.6/7.6 MB 53.5 MB/s eta 0:00:00
_____ 264.6/264.6 KB 25.7 MB/s eta 0:00:00
_____ 158.8/158.8 KB 17.0 MB/s eta 0:00:00
_____ 114.2/114.2 KB 10.3 MB/s eta 0:00:00
_____ 1.6/1.6 MB 74.1 MB/s eta 0:00:00

Mounted at /content/gdrive

#hide
from fastbook import *
```

Other Computer Vision Problems

Multi-Label Classification

The Data

```
from fastai.vision.all import *
path = untar_data(URLs.PASCAL_2007)

100.00% [1637801984/1637796771 02:04<00:00]

df = pd.read_csv(path/'train.csv')
df.head()
```

	fname	labels	is_valid
0	000005.jpg	chair	True
1	000007.jpg	car	True
2	000009.jpg	horse person	True
3	000012.jpg	car	False
4	000016.jpg	bicycle	True

Sidebar: Pandas and DataFrames

```
df.iloc[:,0]

0      000005.jpg
1      000007.jpg
2      000009.jpg
3      000012.jpg
4      000016.jpg
...
5006    009954.jpg
5007    009955.jpg
5008    009958.jpg
5009    009959.jpg
5010    009961.jpg
Name: fname, Length: 5011, dtype: object

df.iloc[0,:]
# Trailing :s are always optional (in numpy, pytorch, pandas, etc.),
# so this is equivalent:
df.iloc[0]
```

```
fname      000005.jpg
labels     chair
is_valid    True
Name: 0, dtype: object

df['fname']

0      000005.jpg
1      000007.jpg
2      000009.jpg
3      000012.jpg
4      000016.jpg
...
5006    009954.jpg
5007    009955.jpg
5008    009958.jpg
5009    009959.jpg
5010    009961.jpg
Name: fname, Length: 5011, dtype: object

tmp_df = pd.DataFrame({'a':[1,2], 'b':[3,4]})
tmp_df
```

	a	b
0	1	3
1	2	4

```
tmp_df['c'] = tmp_df['a']+tmp_df['b']
tmp_df
```

	a	b	c
0	1	3	4
1	2	4	6

End sidebar

▼ Constructing a DataBlock

```
dblock = DataBlock()

dssets = dblock.datasets(df)

len(dssets.train),len(dssets.valid)

(4009, 1002)

x,y = dssets.train[0]
x,y

(fname      008663.jpg
 labels     car person
 is_valid    False
 Name: 4346, dtype: object, fname      008663.jpg
 labels     car person
 is_valid    False
 Name: 4346, dtype: object)

x['fname']

'008663.jpg'

dblock = DataBlock(get_x = lambda r: r['fname'], get_y = lambda r: r['labels'])
dssets = dblock.datasets(df)
dssets.train[0]

('005620.jpg', 'aeroplane')

def get_x(r): return r['fname']
def get_y(r): return r['labels']
dblock = DataBlock(get_x = get_x, get_y = get_y)
```

```

dsets = dblock.datasets(df)
dsets.train[0]

('002549.jpg', 'tvmonitor')

def get_x(r): return path/'train'/r['fname']
def get_y(r): return r['labels'].split(' ')
dblock = DataBlock(get_x = get_x, get_y = get_y)
dsets = dblock.datasets(df)
dsets.train[0]

(Path('/root/.fastai/data/pascal_2007/train/002844.jpg'), ['train'])

dblock = DataBlock(blocks=(ImageBlock, MultiCategoryBlock),
                    get_x = get_x, get_y = get_y)
dsets = dblock.datasets(df)
dsets.train[0]

(PILImage mode=RGB size=500x375,
 TensorMultiCategory([0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]))

idxs = torch.where(dsets.train[0][1]==1.)[0]
dsets.train.vocab[idxs]

(#1) ['dog']

def splitter(df):
    train = df.index[~df['is_valid']].tolist()
    valid = df.index[df['is_valid']].tolist()
    return train,valid

dblock = DataBlock(blocks=(ImageBlock, MultiCategoryBlock),
                    splitter=splitter,
                    get_x=get_x,
                    get_y=get_y)

dsets = dblock.datasets(df)
dsets.train[0]

(PILImage mode=RGB size=500x333,
 TensorMultiCategory([0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.]))

dblock = DataBlock(blocks=(ImageBlock, MultiCategoryBlock),
                    splitter=splitter,
                    get_x=get_x,
                    get_y=get_y,
                    item_tfms = RandomResizedCrop(128, min_scale=0.35))
dls = dblock.dataloaders(df)

dls.show_batch(nrows=1, ncols=3)

```



▼ Binary Cross-Entropy

```

learn = vision_learner(dls, resnet18)

/usr/local/lib/python3.9/dist-packages/torchvision/models/_utils.py:208: UserWarni
warnings.warn(
/usr/local/lib/python3.9/dist-packages/torchvision/models/_utils.py:223: UserWarni
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/
100% 44.7M/44.7M [00:00<00:00, 125MB/s]

x,y = to_cpu(dls.train.one_batch())
activs = learn.model(x)

```

```

activs.shape

      torch.Size([64, 20])

activs[0]

TensorBase([ 0.7476, -1.1988,  4.5421, -1.5915, -0.6749,  0.0343, -2.4930, -0.8330, -0.3817, -1.4876, -0.1683,  2.1547, -3.4151,
-1.1743,  0.1530, -1.6801, -2.3067,  0.7063, -1.3358, -0.3715],
      grad_fn=<AliasBackward0>)

def binary_cross_entropy(inputs, targets):
    inputs = inputs.sigmoid()
    return -torch.where(targets==1, inputs, 1-inputs).log().mean()

loss_func = nn.BCEWithLogitsLoss()
loss = loss_func(activs, y)
loss

TensorMultiCategory(1.0342, grad_fn=<AliasBackward0>)

def say_hello(name, say_what="Hello"): return f"{say_what} {name}."
say_hello('Jeremy'),say_hello('Jeremy', 'Ahoy!')

('Hello Jeremy.', 'Ahoy! Jeremy.')

f = partial(say_hello, say_what="Bonjour")
f("Jeremy"),f("Sylvain")

('Bonjour Jeremy.', 'Bonjour Sylvain.')

learn = vision_learner(dls, resnet50, metrics=partial(accuracy_multi, thresh=0.2))
learn.fine_tune(3, base_lr=3e-3, freeze_epochs=4)

/usr/local/lib/python3.9/dist-packages/torchvision/models/_utils.py:223: UserWarning:
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/
100%          97.8M/97.8M [00:00<00:00, 163MB/s]

epoch  train_loss  valid_loss  accuracy_multi  time
-----
0      0.943426    0.692230      0.235896  00:45
1      0.823277    0.564228      0.285199  00:45
2      0.604020    0.199862      0.827908  00:34
3      0.359526    0.124002      0.944323  00:34
epoch  train_loss  valid_loss  accuracy_multi  time
-----
0      0.131472    0.116906      0.944203  00:38
1      0.116399    0.106551      0.951096  00:35
2      0.096168    0.104706      0.951116  00:35

```



```

learn.metrics = partial(accuracy_multi, thresh=0.1)
learn.validate()

(#2) [0.10470623522996902,0.9281672835350037]

learn.metrics = partial(accuracy_multi, thresh=0.99)
learn.validate()

(#2) [0.10470623522996902,0.9431872963905334]


preds,targs = learn.get_preds()

accuracy_multi(preds, targs, thresh=0.9, sigmoid=False)

TensorBase(0.9571)

xs = torch.linspace(0.05,0.95,29)
accs = [accuracy_multi(preds, targs, thresh=i, sigmoid=False) for i in xs]
plt.plot(xs,accs);

```



▼ Regression

▼ Assemble the Data

```
path = untar_data(URLs.BIWI_HEAD_POSE)
```

100.00% [452321280/452316199 00:35<00:00]

```
#hide
Path.BASE_PATH = path
```

```
path.ls().sorted()
```

```
(#50)
[Path('01'),Path('01.obj'),Path('02'),Path('02.obj'),Path('03'),Path('03.obj'),Path('04'),Path('04.obj'),Path('05'),Path('05.obj')]
```

```
(path/'01').ls().sorted()
```

```
(#1000)
[Path('01/depth.cal'),Path('01/frame_00003_pose.txt'),Path('01/frame_00003_rgb.jpg'),Path('01/frame_00004_pose.txt'),Path('01/frame
```

```
img_files = get_image_files(path)
def img2pose(x): return Path(f'{str(x)[-7]}pose.txt')
img2pose(img_files[0])
```

```
Path('05/frame_00566_pose.txt')
```

```
im = PILImage.create(img_files[0])
im.shape
```

```
(480, 640)
```

```
im.to_thumb(160)
```



```
cal = np.genfromtxt(path/'01'/'rgb.cal', skip_footer=6)
def get_ctr(f):
    ctr = np.genfromtxt(img2pose(f), skip_header=3)
    c1 = ctr[0] * cal[0][0]/ctr[2] + cal[0][2]
    c2 = ctr[1] * cal[1][1]/ctr[2] + cal[1][2]
    return tensor([c1,c2])
```

```
get_ctr(img_files[0])
```

```
tensor([346.0781, 267.9450])
```

```
biwi = DataBlock(
    blocks=(ImageBlock, PointBlock),
    get_items=get_image_files,
    get_y=get_ctr,
```

```


splitter=FuncSplitter(lambda o: o.parent.name=='13'),
batch_tfms=aug_transforms(size=(240,320)),
)

```

```

dls = biwi.dataloaders(path)
dls.show_batch(max_n=9, figsize=(8,6))

```



```

xb,yb = dls.one_batch()
xb.shape,yb.shape

(torch.Size([64, 3, 240, 320]), torch.Size([64, 1, 2]))

yb[0]

TensorPoint([[ 0.1568, -0.0116]], device='cuda:0')

```

▼ Training a Model


```

learn = vision_learner(dls, resnet18, y_range=(-1,1))

def sigmoid_range(x, lo, hi): return torch.sigmoid(x) * (hi-lo) + lo

plot_function(partial(sigmoid_range,lo=-1,hi=1), min=-4, max=4)

```



```

dls.loss_func

FlattenedLoss of MSELoss()

learn.lr_find()

```

```
SuggestedLRs(valley=0.0010000000474974513)
```



```
lr = 1e-2
```

```
learn.fine_tune(3, lr)
```

epoch	train_loss	valid_loss	time
0	0.048756	0.013748	02:16


epoch	train_loss	valid_loss	time
0	0.008029	0.002170	02:22
1	0.003059	0.000180	02:23
2	0.001466	0.000081	02:24

```
math.sqrt(0.0001)
```

```
0.01
```

```
learn.show_results(ds_idx=1, nrows=3, figsize=(6,8))
```

Target/Prediction



Conclusion

▼ Questionnaire

1. How could multi-label classification improve the usability of the bear classifier?
2. How do we encode the dependent variable in a multi-label classification problem?
3. How do you access the rows and columns of a DataFrame as if it was a matrix?
4. How do you get a column by name from a DataFrame?
5. What is the difference between a Dataset and DataLoader?
6. What does a Datasets object normally contain?
7. What does a DataLoaders object normally contain?
8. What does lambda do in Python?
9. What are the methods to customize how the independent and dependent variables are created with the data block API?
10. Why is softmax not an appropriate output activation function when using a one hot encoded target?
11. Why is nll_loss not an appropriate loss function when using a one-hot-encoded target?
12. What is the difference between nn.BCELoss and nn.BCEWithLogitsLoss?
13. Why can't we use regular accuracy in a multi-label problem?
14. When is it okay to tune a hyperparameter on the validation set?
15. How is y_range implemented in fastai? (See if you can implement it yourself and test it without peeking!)
16. What is a regression problem? What loss function should you use for such a problem?

17. What do you need to do to make sure the `fastai` library applies the same data augmentation to your input images and your target point coordinates?

▼ Further Research

1. Read a tutorial about Pandas DataFrames and experiment with a few methods that look interesting to you. See the book's website for recommended tutorials.
2. Retrain the bear classifier using multi-label classification. See if you can make it work effectively with images that don't contain any bears, including showing that information in the web application. Try an image with two different kinds of bears. Check whether the accuracy on the single-label dataset is impacted using multi-label classification.

